

the transformer architecture is described here in enough detail with several lines of ordinary prose so that the extractor has something substantial to read back out of the text layer, which is what tells a real paper apart from a scanned page that only yields a handful of stray glyphs and no sentences whatsoever. Attention weights are computed over the whole input sequence.